

```
In [1]: import numpy as np
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.utils import to_categorical
```

```
In [2]: (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [3]: # Reshape to (28, 28, 1) for CNN input and normalize the pixel values (0-255)
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1).astype('float32') / 255
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1).astype('float32') / 255.0
```

```
In [4]: # One-hot encode the labels
y_train = to_categorical(y_train, num_classes=10)
y_test = to_categorical(y_test, num_classes=10)
```

```
In [5]: # Build the CNN model
model = Sequential()

# First convolutional layer + pooling
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(MaxPooling2D(pool_size=(2, 2)))

# Second convolutional layer + pooling
model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
In [6]: # Flatten the output for the fully connected layers
model.add(Flatten())

# Fully connected layer with 128 neurons
model.add(Dense(128, activation='relu'))

# Dropout layer to prevent overfitting
model.add(Dropout(0.5))

# Output layer with 10 neurons for the 10 digit classes (0-9)
model.add(Dense(10, activation='softmax'))
```

```
In [7]: # Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['acc

# Train the model
history = model.fit(X_train, y_train, epochs=10, batch_size=128, validation_sp

# Evaluate the model on the test set
test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=2)
```

Epoch 1/10

375/375 - 39s - loss: 0.3561 - accuracy: 0.8916 - val\_loss: 0.0798 - val\_accu  
racy: 0.9768 - 39s/epoch - 105ms/step

Epoch 2/10

375/375 - 36s - loss: 0.1113 - accuracy: 0.9675 - val\_loss: 0.0566 - val\_accu  
racy: 0.9848 - 36s/epoch - 97ms/step

Epoch 3/10

375/375 - 36s - loss: 0.0809 - accuracy: 0.9756 - val\_loss: 0.0450 - val\_accu  
racy: 0.9876 - 36s/epoch - 96ms/step

Epoch 4/10

375/375 - 36s - loss: 0.0648 - accuracy: 0.9806 - val\_loss: 0.0394 - val\_accu  
racy: 0.9883 - 36s/epoch - 95ms/step

Epoch 5/10

375/375 - 36s - loss: 0.0554 - accuracy: 0.9836 - val\_loss: 0.0370 - val\_accu  
racy: 0.9895 - 36s/epoch - 95ms/step

Epoch 6/10

375/375 - 35s - loss: 0.0470 - accuracy: 0.9859 - val\_loss: 0.0401 - val\_accu  
racy: 0.9882 - 35s/epoch - 94ms/step

Epoch 7/10

375/375 - 35s - loss: 0.0418 - accuracy: 0.9871 - val\_loss: 0.0387 - val\_accu  
racy: 0.9892 - 35s/epoch - 93ms/step

Epoch 8/10

375/375 - 36s - loss: 0.0370 - accuracy: 0.9887 - val\_loss: 0.0348 - val\_accu  
racy: 0.9904 - 36s/epoch - 96ms/step

Epoch 9/10

375/375 - 35s - loss: 0.0326 - accuracy: 0.9900 - val\_loss: 0.0343 - val\_accu  
racy: 0.9913 - 35s/epoch - 94ms/step

Epoch 10/10

375/375 - 34s - loss: 0.0300 - accuracy: 0.9903 - val\_loss: 0.0396 - val\_accu  
racy: 0.9898 - 34s/epoch - 92ms/step

313/313 - 3s - loss: 0.0263 - accuracy: 0.9923 - 3s/epoch - 10ms/step

```
In [8]: print(f'Test Accuracy: {test_accuracy:.4f}')
```

Test Accuracy: 0.9923